



Roy's AI Lab
THINK · BUILD · SCALE

FREE GUIDE

50 Q&A

INTERVIEW PREP SERIES

AWS for AI Engineers

50 real, scenario-based interview questions — with model answers, common mistakes, and follow-ups — for GenAI & AI Engineers working with AWS.

50

Q&A Pairs

4

Difficulty Levels

100%

Real Interview Signal

Q07 · SCENARIO-BASED · BEGINNER

INTERVIEWER

"How would you reduce cost when calling Bedrock models at scale in production?"

MODEL ANSWER

I'd start by right-sizing the model per task — not every call needs the largest model. Then layer in prompt caching, batching, and response streaming to cut both latency and token spend.



Ritesh Rai (Roy)

GenAI Engineer · Founder, Roy's AI Lab

2+ years building production AI systems · 500K+ LinkedIn impressions · Top 1% Creator on Topmate

Swipe to start →

01 / 51

What is AWS, and why is it widely used for AI applications?

What is AWS, and why is it widely used for AI applications?

| ANSWER

AWS (Amazon Web Services) is a cloud computing platform that provides on-demand services like computing, storage, databases, networking, security, and AI tools.

Instead of buying and managing physical servers, companies can launch resources on AWS within minutes and pay only for what they use.

For AI applications, AWS offers services to train machine learning models, deploy LLMs, store datasets, run APIs, monitor applications, and automatically scale infrastructure when traffic increases.

For example, an AI chatbot might use S3 to store documents, Bedrock to access an LLM, OpenSearch for vector search, Lambda for backend logic, and CloudWatch for monitoring — all within AWS.

This allows teams to focus on building AI products rather than managing servers.

| WHY INTERVIEWERS ASK THIS

They want to check whether you understand AWS as a complete cloud platform — not just a collection of services.

| COMMON MISTAKES

- Saying AWS is only used for hosting websites.
- Thinking AWS is only for Machine Learning.
- Explaining cloud computing without mentioning scalability or pay-as-you-go pricing.

| FOLLOW-UP QUESTIONS

- Why do AI startups prefer AWS over on-premise servers?
- What AWS services have you personally used?
- What are the main advantages of cloud computing?

PRODUCTION INSIGHT

In production, AWS lets teams build highly available AI systems without maintaining physical infrastructure. Features like auto scaling, managed databases, and serverless services reduce operational work while improving reliability.

What are the main AWS services an AI Engineer should know?

What are the main AWS services an AI Engineer should know?

ANSWER

An AI Engineer doesn't need to know every AWS service, but should be comfortable with the ones commonly used in AI applications. Some of the most important services are:

- EC2 — Runs applications on virtual machines.
- S3 — Stores datasets, models, and documents.
- IAM — Manages users, roles, and permissions.
- Lambda — Runs backend code without managing servers.
- API Gateway — Exposes AI models through APIs.
- Bedrock — Accesses foundation models like Claude and Llama.
- SageMaker — Builds, trains, and deploys ML models.
- OpenSearch — Enables vector search for RAG applications.
- CloudWatch — Monitors logs, metrics, and system health.
- ECR — Stores Docker images.
- ECS/EKS — Deploys containerized AI applications.

Knowing how these services work together is more valuable than memorizing hundreds of AWS services.

WHY INTERVIEWERS ASK THIS

They want to know whether you understand the AWS ecosystem used in real AI projects.

COMMON MISTAKES

- Listing random AWS services without explaining their purpose.
- Forgetting IAM or CloudWatch.
- Ignoring deployment-related services.

FOLLOW-UP QUESTIONS

- Which AWS services did you use in your last project?
- Which service would you use to deploy an LLM?
- Where would you store embeddings?

PRODUCTION INSIGHT

Most production AI systems combine multiple AWS services rather than relying on just one. Understanding how these services interact is a key interview differentiator.

What is Amazon EC2?

What is Amazon EC2?

| ANSWER

Amazon EC2 (Elastic Compute Cloud) provides virtual servers in the cloud.

Instead of buying physical machines, you launch an EC2 instance with the required CPU, RAM, GPU, storage, and operating system.

AI Engineers commonly use EC2 to host FastAPI applications, run LangGraph workflows, deploy RAG systems, serve LLM APIs, or perform model inference.

For GPU-heavy workloads, AWS offers GPU instances that significantly speed up deep learning and inference tasks.

Since EC2 is fully configurable, you are responsible for installing software, updating packages, configuring security, and managing scaling.

| WHY INTERVIEWERS ASK THIS

They want to confirm you understand the difference between managing your own server and using serverless services.

| COMMON MISTAKES

- Saying EC2 is just “a server.”
- Forgetting that you manage the operating system yourself.
- Confusing EC2 with Lambda.

| FOLLOW-UP QUESTIONS

- When would you choose EC2 instead of Lambda?
- Can EC2 run Docker containers?
- What happens if an EC2 instance crashes?

PRODUCTION INSIGHT

Many AI companies use EC2 with Docker and Nginx to deploy production AI APIs because it provides greater flexibility than serverless solutions.

What is Amazon S3, and how is it used in AI projects?

What is Amazon S3, and how is it used in AI projects?

| ANSWER

Amazon S3 (Simple Storage Service) is AWS's object storage service.

It stores files like PDFs, images, datasets, trained models, embeddings, logs, videos, and backups.

In AI projects, S3 is often the first stop in the data pipeline. Documents uploaded by users are stored in S3, then processed to create embeddings before being indexed into a vector database.

S3 is highly durable, scalable, and cost-effective. It automatically replicates data across multiple devices, making it a reliable storage solution.

Since it can store virtually unlimited data, it is commonly used for data lakes and model storage.

| WHY INTERVIEWERS ASK THIS

They want to see if you know where AI systems store large amounts of data.

| COMMON MISTAKES

- Calling S3 a database.
- Thinking S3 is only for images.
- Confusing object storage with block storage.

| FOLLOW-UP QUESTIONS

- Can you directly query data from S3?
- What is an S3 bucket?
- How does S3 differ from EBS?

PRODUCTION INSIGHT

In many production RAG systems, documents are uploaded to S3 first. An event then triggers processing, embedding generation, and indexing automatically.

What is IAM, and why is it important?

What is IAM, and why is it important?

| ANSWER

IAM (Identity and Access Management) controls who can access AWS resources and what actions they are allowed to perform.

Instead of sharing AWS account credentials, IAM lets you create users, groups, and roles with specific permissions.

For example, an AI application may need permission to read files from S3 but should not be allowed to delete them. IAM policies enforce these rules.

Using the principle of least privilege, every user or service gets only the permissions it actually needs, reducing security risks.

IAM is one of the most important AWS services because almost every AWS resource is protected through IAM permissions.

| WHY INTERVIEWERS ASK THIS

Security is critical in production systems, and IAM is the foundation of AWS security.

| COMMON MISTAKES

- Using the root account for daily work.
- Giving administrator access to every service.
- Ignoring least privilege.

| FOLLOW-UP QUESTIONS

- What is an IAM Role?
- What is the difference between a user and a role?
- What happens if an application lacks S3 permissions?

PRODUCTION INSIGHT

Most production AI applications use IAM Roles attached to EC2 instances, Lambda functions, or ECS tasks instead of storing AWS access keys in code.

What is AWS Lambda, and when should an AI Engineer use it?

What is AWS Lambda, and when should an AI Engineer use it?

| ANSWER

AWS Lambda is a serverless compute service that lets you run code without managing servers. You simply upload your code, and AWS automatically executes it whenever an event occurs.

For AI applications, Lambda is commonly used for lightweight backend tasks such as processing uploaded documents, triggering embedding generation, calling an LLM API, sending notifications, or handling API requests.

One of its biggest advantages is that you only pay when your code runs. AWS automatically scales the function based on incoming requests, so you don't have to worry about provisioning servers.

However, Lambda is not ideal for long-running AI inference or GPU-based workloads because it has execution time and memory limits.

| WHY INTERVIEWERS ASK THIS

They want to know if you understand when serverless architecture is the right choice and when it isn't.

| COMMON MISTAKES

- Saying Lambda can replace every EC2 server.
- Ignoring Lambda's execution time limits.
- Using Lambda for GPU-intensive model inference.

| FOLLOW-UP QUESTIONS

- What triggers a Lambda function?
- Can Lambda access S3?
- When would you choose EC2 instead of Lambda?

PRODUCTION INSIGHT

In production RAG systems, uploading a PDF to S3 often triggers a Lambda function that extracts the text, creates embeddings, and stores them in a vector database automatically.

What is API Gateway, and why is it commonly used with Lambda?

What is API Gateway, and why is it commonly used with Lambda?

| ANSWER

API Gateway is a managed AWS service that allows developers to create, publish, and manage APIs.

Instead of exposing a Lambda function directly to users, API Gateway receives the HTTP request, validates it, handles authentication if needed, and then forwards the request to Lambda.

For example, if a user sends a prompt to your chatbot, the request first reaches API Gateway. It then invokes Lambda, which processes the request and returns the response.

API Gateway also provides features like rate limiting, authentication, request validation, and monitoring, making it suitable for production environments.

| WHY INTERVIEWERS ASK THIS

They want to check whether you understand how backend APIs are exposed securely in AWS.

| COMMON MISTAKES

- Thinking API Gateway executes the application logic.
- Confusing API Gateway with a Load Balancer.
- Forgetting that API Gateway can trigger Lambda.

| FOLLOW-UP QUESTIONS

- Can API Gateway call EC2 instead of Lambda?
- What authentication methods does API Gateway support?
- Why not expose Lambda directly?

PRODUCTION INSIGHT

Many AI startups expose their LLM APIs using API Gateway because it provides authentication, throttling, logging, and monitoring without additional infrastructure.

What is a VPC, and why is it important?

What is a VPC, and why is it important?

| ANSWER

A VPC (Virtual Private Cloud) is your own private network inside AWS.

It allows you to control how AWS resources communicate with each other and with the internet.

Inside a VPC, you can create public and private subnets, configure routing, control network traffic, and isolate sensitive resources.

For example, an AI application may place its API server in a public subnet while keeping databases, Redis, and vector databases inside private subnets where they cannot be accessed directly from the internet.

Using a VPC improves security and gives you full control over your network architecture.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand basic AWS networking concepts used in production systems.

| COMMON MISTAKES

- Saying a VPC is just another server.
- Confusing VPC with EC2.
- Assuming every AWS resource is publicly accessible.

| FOLLOW-UP QUESTIONS

- What is a subnet?
- What is the difference between public and private subnets?
- Why keep databases in private subnets?

PRODUCTION INSIGHT

Almost every enterprise AI application runs inside a VPC because it allows secure communication between application servers, databases, caches, and internal services.

What are Security Groups in AWS?

What are Security Groups in AWS?

| ANSWER

Security Groups act like virtual firewalls for AWS resources such as EC2 instances.

They control which inbound and outbound traffic is allowed.

For example, you may allow HTTP traffic on port 80, HTTPS on port 443, and SSH access on port 22 only from your office IP address.

If traffic doesn't match an allowed rule, AWS blocks it automatically.

Unlike traditional firewalls, Security Groups are stateful. This means if an incoming request is allowed, the response is automatically allowed without creating another rule.

| WHY INTERVIEWERS ASK THIS

They want to evaluate whether you understand one of the most basic security features in AWS.

| COMMON MISTAKES

- Opening every port to the internet.
- Confusing Security Groups with IAM.
- Not restricting SSH access.

| FOLLOW-UP QUESTIONS

- Are Security Groups stateful or stateless?
- Can multiple Security Groups be attached to one EC2 instance?
- What is the difference between Security Groups and Network ACLs?

PRODUCTION INSIGHT

One of the most common production mistakes is allowing unrestricted access (0.0.0.0/0) to sensitive ports like SSH. Following the principle of least privilege helps prevent unnecessary security risks.

What is Amazon ECR, and why is it used with Docker?

What is Amazon ECR, and why is it used with Docker?

| ANSWER

Amazon ECR (Elastic Container Registry) is AWS's managed service for storing Docker container images.

When an AI application is packaged into a Docker image, that image needs to be stored somewhere before deployment. ECR provides a secure and scalable place to store these images.

Services like ECS, EKS, and even some CI/CD pipelines pull Docker images directly from ECR during deployment.

Using ECR also simplifies version management, integrates with IAM for secure access, and works seamlessly with other AWS services.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how containerized AI applications are deployed in AWS.

| COMMON MISTAKES

- Thinking ECR runs Docker containers.
- Confusing ECR with ECS.
- Assuming Docker images are stored on EC2 permanently.

| FOLLOW-UP QUESTIONS

- What is the difference between ECR and ECS?
- Can Kubernetes pull images from ECR?
- How do you update a Docker image stored in ECR?

PRODUCTION INSIGHT

A common production workflow is: build the Docker image → push it to ECR → deploy it using ECS or EKS. This creates a clean and repeatable deployment process for AI applications.

What is Amazon ECS, and when should an AI Engineer use it?

What is Amazon ECS, and when should an AI Engineer use it?

| ANSWER

Amazon ECS (Elastic Container Service) is AWS's managed service for running Docker containers without managing the underlying container orchestration yourself.

Instead of manually starting Docker containers on multiple EC2 instances, ECS handles scheduling, deployment, scaling, and monitoring for you.

AI Engineers commonly use ECS to deploy FastAPI applications, RAG systems, LangGraph workflows, and LLM inference services packaged as Docker containers.

ECS works well when your application needs long-running services and predictable performance but doesn't require the complexity of Kubernetes.

| WHY INTERVIEWERS ASK THIS

They want to know if you understand modern deployment practices for AI applications.

| COMMON MISTAKES

- Thinking ECS is the same as ECR.
- Assuming ECS stores Docker images.
- Confusing ECS with Kubernetes.

| FOLLOW-UP QUESTIONS

- What is the difference between ECS and EKS?
- Can ECS use EC2 instances?
- What is AWS Fargate?

PRODUCTION INSIGHT

Many startups choose ECS because it's simpler to manage than Kubernetes while still providing reliable container deployment and auto scaling.

What is Amazon EKS, and why do large companies use it?

What is Amazon EKS, and why do large companies use it?

| ANSWER

Amazon EKS (Elastic Kubernetes Service) is AWS's managed Kubernetes service.

It allows companies to deploy and manage containerized applications using Kubernetes without managing the Kubernetes control plane themselves.

Large organizations prefer EKS because Kubernetes is cloud-independent, highly scalable, and widely adopted across the industry.

For AI applications, EKS is commonly used to deploy multiple microservices such as APIs, embedding services, vector databases, inference servers, and monitoring tools.

Although EKS offers greater flexibility than ECS, it is also more complex to configure and maintain.

| WHY INTERVIEWERS ASK THIS

They want to see whether you understand the difference between simple container deployment and enterprise-scale orchestration.

| COMMON MISTAKES

- Saying EKS is just another Docker service.
- Thinking Kubernetes is only for large companies.
- Confusing EKS with ECS.

| FOLLOW-UP QUESTIONS

- When would you choose EKS over ECS?
- Does EKS require Docker?
- What is a Kubernetes Pod?

PRODUCTION INSIGHT

Large AI platforms often standardize on Kubernetes because it provides portability across AWS, Azure, and GCP while supporting advanced deployment strategies.

What is AWS Fargate?

What is AWS Fargate?

| ANSWER

AWS Fargate is a serverless compute engine for containers.

Normally, when running containers, you must manage EC2 instances. With Fargate, AWS manages the infrastructure for you.

You simply provide your Docker container, specify the required CPU and memory, and AWS launches and scales it automatically.

This makes Fargate a great choice for AI APIs, document processing services, and backend applications that don't require GPU support.

It reduces operational work because there are no servers to patch, monitor, or maintain.

| WHY INTERVIEWERS ASK THIS

They want to know whether you're familiar with serverless container deployment.

| COMMON MISTAKES

- Thinking Fargate is a replacement for Docker.
- Assuming Fargate supports every GPU workload.
- Confusing Fargate with Lambda.

| FOLLOW-UP QUESTIONS

- What is the difference between ECS on EC2 and ECS on Fargate?
- Is Fargate serverless?
- When should you avoid using Fargate?

PRODUCTION INSIGHT

Many production AI APIs with moderate traffic run on ECS with Fargate because it combines the flexibility of containers with the simplicity of serverless infrastructure.

What is Amazon SageMaker?

What is Amazon SageMaker?

| ANSWER

Amazon SageMaker is AWS's managed machine learning platform.

It helps data scientists and AI Engineers build, train, tune, deploy, and monitor machine learning models from a single platform.

Instead of manually setting up GPU servers and ML environments, SageMaker provides managed notebooks, training jobs, model endpoints, and deployment tools.

It supports popular frameworks like TensorFlow, PyTorch, XGBoost, and Scikit-learn.

For traditional machine learning projects, SageMaker significantly reduces infrastructure management and speeds up development.

| WHY INTERVIEWERS ASK THIS

They want to know whether you've worked with AWS's primary machine learning service.

| COMMON MISTAKES

- Thinking SageMaker is only for training models.
- Confusing SageMaker with Bedrock.
- Assuming SageMaker automatically creates ML models.

| FOLLOW-UP QUESTIONS

- What is a SageMaker Endpoint?
- Can SageMaker deploy PyTorch models?
- What is SageMaker Studio?

PRODUCTION INSIGHT

Many enterprises train their ML models using SageMaker and then deploy them as managed endpoints for real-time predictions with built-in monitoring and scaling.

What is Amazon Bedrock, and why is it important for Generative AI?

What is Amazon Bedrock, and why is it important for Generative AI?

| ANSWER

Amazon Bedrock is AWS's fully managed Generative AI service that provides access to foundation models through a single API.

Instead of hosting your own LLM, you can directly use models from providers like Anthropic, Meta, Amazon, Mistral, and others without managing infrastructure.

For AI Engineers, Bedrock simplifies tasks like chatbot development, document summarization, content generation, question answering, and RAG applications.

AWS handles model hosting, scaling, availability, and infrastructure, allowing developers to focus on building AI features.

Bedrock is especially useful when you want production-ready access to multiple LLMs without maintaining GPU servers.

| WHY INTERVIEWERS ASK THIS

Generative AI interviews increasingly include Bedrock because it's becoming one of the most widely used managed LLM platforms on AWS.

| COMMON MISTAKES

- Thinking Bedrock trains foundation models.
- Confusing Bedrock with SageMaker.
- Assuming Bedrock only supports Amazon's own models.

| FOLLOW-UP QUESTIONS

- Which foundation models are available in Bedrock?
- When would you choose Bedrock over SageMaker?
- Can Bedrock be used in a RAG application?

PRODUCTION INSIGHT

Many production GenAI applications use Bedrock because it provides secure, scalable access to multiple LLM providers through a single API, reducing infrastructure complexity and accelerating development.

What is Amazon OpenSearch, and how is it used in RAG applications?

What is Amazon OpenSearch, and how is it used in RAG applications?

| ANSWER

Amazon OpenSearch is a managed search and analytics service that can also be used as a vector database for Retrieval-Augmented Generation (RAG) applications.

In a RAG pipeline, documents are first converted into embeddings using an embedding model. These embeddings are then stored in OpenSearch. When a user asks a question, the query is also converted into an embedding, and OpenSearch finds the most similar documents using vector similarity search.

Besides vector search, OpenSearch also supports keyword search, filtering, logging, and analytics, making it useful for enterprise AI applications.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand where embeddings are stored and how retrieval works in production AI systems.

| COMMON MISTAKES

- Saying OpenSearch is only used for logs.
- Confusing OpenSearch with Amazon S3.
- Thinking OpenSearch generates embeddings.

| FOLLOW-UP QUESTIONS

- What is vector search?
- Can OpenSearch perform hybrid search?
- What is the difference between OpenSearch and Pinecone?

PRODUCTION INSIGHT

Many enterprise RAG systems combine vector search with keyword search in OpenSearch to improve retrieval accuracy instead of relying only on semantic similarity.

What is Amazon DynamoDB, and when should an AI Engineer use it?

What is Amazon DynamoDB, and when should an AI Engineer use it?

| ANSWER

Amazon DynamoDB is AWS's fully managed NoSQL database designed for high performance and low latency.

Unlike relational databases, DynamoDB stores data as key-value pairs or documents, making it suitable for applications that need fast reads and writes.

AI Engineers commonly use DynamoDB to store chat history, user sessions, prompt logs, API metadata, feedback, and application state.

Since AWS automatically handles scaling, backups, and availability, developers don't need to manage database servers.

| WHY INTERVIEWERS ASK THIS

They want to check whether you understand where AI applications store structured application data.

| COMMON MISTAKES

- Calling DynamoDB a SQL database.
- Using DynamoDB for complex joins.
- Confusing DynamoDB with vector databases.

| FOLLOW-UP QUESTIONS

- What is the difference between DynamoDB and RDS?
- Can DynamoDB store JSON?
- Why is DynamoDB considered serverless?

PRODUCTION INSIGHT

Many chatbot applications store conversation history in DynamoDB because it offers low latency, automatic scaling, and high availability.

What is Amazon RDS, and when should you choose it over DynamoDB?

What is Amazon RDS, and when should you choose it over DynamoDB?

| ANSWER

Amazon RDS (Relational Database Service) is AWS's managed service for relational databases such as PostgreSQL, MySQL, MariaDB, Oracle, and SQL Server.

Choose RDS when your application requires SQL queries, relationships between tables, transactions, and data consistency.

For example, an AI platform may store users, subscriptions, billing information, and application metadata inside PostgreSQL running on RDS.

Unlike DynamoDB, RDS supports joins, foreign keys, indexes, and complex SQL queries.

AWS manages backups, software updates, monitoring, and failover, reducing operational work.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand when to choose SQL databases instead of NoSQL databases.

| COMMON MISTAKES

- Saying RDS is only for storing AI models.
- Thinking RDS scales exactly like DynamoDB.
- Confusing RDS with EC2-hosted databases.

| FOLLOW-UP QUESTIONS

- When would you choose PostgreSQL over DynamoDB?
- What databases does RDS support?
- What is Multi-AZ deployment?

PRODUCTION INSIGHT

Many AI products use PostgreSQL on RDS for transactional data while using OpenSearch or Pinecone separately for vector search.

What is Amazon CloudWatch, and why is it important?

What is Amazon CloudWatch, and why is it important?

| ANSWER

Amazon CloudWatch is AWS's monitoring and observability service.

It collects logs, metrics, and performance data from AWS resources like EC2, Lambda, ECS, API Gateway, and RDS.

For AI Engineers, CloudWatch helps monitor API response times, model latency, CPU usage, memory utilization, application logs, and error rates.

You can also configure alarms that notify your team whenever a resource crosses a predefined threshold, helping detect issues before users are affected.

Monitoring is just as important as deployment because production systems must be continuously observed.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how production applications are monitored after deployment.

| COMMON MISTAKES

- Thinking CloudWatch is only for viewing logs.
- Ignoring monitoring completely.
- Assuming logging and monitoring are the same thing.

| FOLLOW-UP QUESTIONS

- Can CloudWatch trigger alerts?
- What is the difference between logs and metrics?
- Can Lambda logs be viewed in CloudWatch?

PRODUCTION INSIGHT

One of the first things production teams investigate during an outage is CloudWatch, because it provides visibility into application health, resource utilization, and failures.

What is Amazon SQS, and why is it useful in AI applications?

What is Amazon SQS, and why is it useful in AI applications?

| ANSWER

Amazon SQS (Simple Queue Service) is a fully managed message queue that helps different parts of an application communicate asynchronously.

Instead of processing everything immediately, one service can place a message into an SQS queue while another service processes it later.

For example, when a user uploads 1,000 PDFs, the upload service can place processing requests into SQS. Multiple worker services then read messages from the queue and generate embeddings in parallel.

Using SQS improves scalability, prevents system overload, and makes applications more reliable during traffic spikes.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand asynchronous architecture used in production AI systems.

| COMMON MISTAKES

- Thinking SQS stores permanent data.
- Using SQS as a database.
- Assuming SQS processes messages itself.

| FOLLOW-UP QUESTIONS

- What is the difference between SQS and SNS?
- What happens if a worker crashes while processing a message?
- What is a Dead Letter Queue (DLQ)?

PRODUCTION INSIGHT

Large-scale document processing systems almost always use message queues like SQS to decouple services, improve reliability, and process thousands of AI tasks concurrently.

What is Amazon SNS, and how is it different from Amazon SQS?

What is Amazon SNS, and how is it different from Amazon SQS?

| ANSWER

Amazon SNS (Simple Notification Service) is a fully managed messaging service used to send notifications to multiple subscribers at the same time.

Unlike SQS, which stores messages in a queue until they are processed, SNS immediately pushes messages to its subscribers such as Lambda functions, email, SMS, HTTP endpoints, or SQS queues.

For example, after an AI model finishes processing a document, SNS can notify users by email, trigger another Lambda function, and send an event to multiple downstream services — all at the same time.

Think of SNS as a broadcasting service, while SQS acts as a waiting queue for background processing.

| WHY INTERVIEWERS ASK THIS

They want to check whether you understand event-driven architectures and the difference between queues and notifications.

| COMMON MISTAKES

- Saying SNS stores messages permanently.
- Confusing SNS with SQS.
- Assuming SNS guarantees message processing.

| FOLLOW-UP QUESTIONS

- Can SNS send messages to SQS?
- What is the Pub/Sub model?
- When would you use SNS instead of SQS?

PRODUCTION INSIGHT

Many production AI systems use SNS and SQS together — SNS broadcasts an event, while multiple SQS queues receive and process it independently.

What is AWS EventBridge?

What is AWS EventBridge?

| ANSWER

AWS EventBridge is an event bus that connects different AWS services and applications without requiring direct integration between them.

Instead of one service calling another directly, services publish events to EventBridge. Based on predefined rules, EventBridge routes those events to the appropriate destination.

For example, when a document is uploaded to S3, EventBridge can trigger a Lambda function, start a Step Functions workflow, or send a notification — all automatically.

This creates loosely coupled applications that are easier to maintain and scale.

| WHY INTERVIEWERS ASK THIS

They want to know if you're familiar with modern event-driven architectures used in production systems.

| COMMON MISTAKES

- Confusing EventBridge with SQS.
- Assuming EventBridge stores data.
- Thinking EventBridge replaces Lambda.

| FOLLOW-UP QUESTIONS

- What events can EventBridge capture?
- How is EventBridge different from SNS?
- Can EventBridge trigger Step Functions?

PRODUCTION INSIGHT

Large AI platforms often use EventBridge to coordinate multiple services because it reduces dependencies between applications and simplifies workflows.

What are AWS Step Functions, and why are they useful for AI workflows?

What are AWS Step Functions, and why are they useful for AI workflows?

| ANSWER

AWS Step Functions allow you to build workflows by connecting multiple AWS services into a sequence of steps.

Instead of writing one large application that performs every task, you divide the workflow into smaller steps.

For example, in a RAG pipeline, the workflow might be: Upload PDF → Extract Text → Generate Embeddings → Store in Vector Database → Notify User.

If one step fails, Step Functions can automatically retry or handle the failure without restarting the entire workflow.

This makes complex AI pipelines easier to manage and monitor.

| WHY INTERVIEWERS ASK THIS

They want to see if you understand workflow orchestration in production AI systems.

| COMMON MISTAKES

- Treating Step Functions like Lambda.
- Building long workflows inside one Lambda function.
- Ignoring retry and error-handling features.

| FOLLOW-UP QUESTIONS

- Can Step Functions call Lambda?
- What happens if one step fails?
- When should you use Step Functions?

PRODUCTION INSIGHT

Many enterprise AI pipelines use Step Functions because they provide visual workflows, automatic retries, and better monitoring than manually chaining multiple services.

What is Amazon ElastiCache (Redis), and why is it used in AI applications?

What is Amazon ElastiCache (Redis), and why is it used in AI applications?

| ANSWER

Amazon ElastiCache is AWS's managed in-memory caching service that supports Redis and Memcached.

Redis stores frequently accessed data in memory, making it much faster than reading from a database every time.

AI Engineers commonly use Redis to store user sessions, conversation history, authentication tokens, API responses, temporary embeddings, and application cache.

Using Redis reduces database load and improves response times, especially for chatbot and LLM applications that require low latency.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand caching and performance optimization.

| COMMON MISTAKES

- Using Redis as permanent storage.
- Assuming Redis replaces a database.
- Forgetting that Redis stores data primarily in memory.

| FOLLOW-UP QUESTIONS

- Why is Redis faster than a database?
- Can Redis be used for chat history?
- What is cache invalidation?

PRODUCTION INSIGHT

Many production chatbot systems cache frequently asked questions and user session data in Redis, reducing API latency and improving user experience.

What is AWS Auto Scaling, and why is it important for AI applications?

What is AWS Auto Scaling, and why is it important for AI applications?

| ANSWER

AWS Auto Scaling automatically increases or decreases computing resources based on application demand.

When traffic increases, AWS launches additional instances or containers. When traffic decreases, AWS removes unused resources to reduce costs.

For example, if an AI chatbot suddenly receives thousands of requests during a product launch, Auto Scaling can automatically add more servers to handle the increased load without manual intervention.

This helps maintain application performance while avoiding unnecessary infrastructure costs during low-traffic periods.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how production systems handle changing workloads.

| COMMON MISTAKES

- Thinking Auto Scaling only adds servers.
- Assuming scaling happens instantly.
- Ignoring cost optimization benefits.

| FOLLOW-UP QUESTIONS

- What metrics can trigger Auto Scaling?
- Can ECS services be auto-scaled?
- How does Auto Scaling reduce costs?

PRODUCTION INSIGHT

Production AI systems are designed for unpredictable traffic. Auto Scaling ensures users continue receiving fast responses during peak demand while minimizing infrastructure costs during quieter periods.

What is an Application Load Balancer (ALB), and why is it important?

What is an Application Load Balancer (ALB), and why is it important?

| ANSWER

An Application Load Balancer (ALB) distributes incoming HTTP and HTTPS requests across multiple servers or containers.

Instead of sending all user requests to a single server, the ALB intelligently routes traffic to healthy instances. This improves availability, scalability, and fault tolerance.

For example, if your AI chatbot runs on four EC2 instances, the ALB ensures incoming requests are shared among all four. If one instance crashes, the ALB automatically stops sending traffic to it.

This prevents downtime and provides a better experience for users.

| WHY INTERVIEWERS ASK THIS

They want to know if you understand how production applications handle high traffic and server failures.

| COMMON MISTAKES

- Thinking the Load Balancer stores application data.
- Assuming it increases server performance directly.
- Confusing Load Balancer with API Gateway.

| FOLLOW-UP QUESTIONS

- What is the difference between ALB and NLB?
- Can ECS services use an ALB?
- How does the Load Balancer detect unhealthy instances?

PRODUCTION INSIGHT

Almost every production AI API is placed behind a Load Balancer because it improves reliability, enables zero-downtime deployments, and supports automatic scaling.

What is Amazon Route 53?

What is Amazon Route 53?

| ANSWER

Amazon Route 53 is AWS's managed Domain Name System (DNS) service.

Its main job is to translate human-readable domain names like `api.company.com` into the IP addresses of your servers.

Without Route 53, users would need to remember IP addresses instead of domain names.

Route 53 also supports health checks, traffic routing, failover, and latency-based routing, making applications more reliable.

For example, when someone opens your AI application, Route 53 directs them to the correct Load Balancer or server.

| WHY INTERVIEWERS ASK THIS

They want to check your understanding of how users reach applications hosted on AWS.

| COMMON MISTAKES

- Thinking Route 53 hosts websites.
- Confusing DNS with web hosting.
- Assuming Route 53 is only for buying domains.

| FOLLOW-UP QUESTIONS

- What is DNS?
- Can Route 53 perform health checks?
- What happens if one AWS region goes down?

PRODUCTION INSIGHT

Many global AI applications use Route 53 with health checks to automatically redirect users to healthy regions if one region becomes unavailable.

What is Amazon CloudFront, and why is it used?

What is Amazon CloudFront, and why is it used?

| ANSWER

Amazon CloudFront is AWS's Content Delivery Network (CDN).

It improves application performance by caching content at edge locations around the world.

Instead of downloading files from the original server every time, users receive cached content from the nearest edge location, reducing latency.

CloudFront is commonly used to deliver images, JavaScript files, CSS files, videos, documentation, and static web pages.

Although it is mainly designed for static content, it can also improve API performance in certain scenarios.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how global applications improve speed and reduce latency.

| COMMON MISTAKES

- Thinking CloudFront stores permanent data.
- Confusing CloudFront with S3.
- Assuming CloudFront only works for websites.

| FOLLOW-UP QUESTIONS

- What is a CDN?
- Can CloudFront work with S3?
- How does CloudFront reduce latency?

PRODUCTION INSIGHT

Many AI applications store frontend assets in S3 and deliver them through CloudFront, allowing users worldwide to experience faster page load times.

What is AWS Secrets Manager, and why should you use it?

What is AWS Secrets Manager, and why should you use it?

| ANSWER

AWS Secrets Manager is a service for securely storing sensitive information such as API keys, database passwords, authentication tokens, and third-party credentials.

Instead of hardcoding secrets into your application or storing them in configuration files, your application retrieves them securely from Secrets Manager whenever needed.

This improves security because secrets can be encrypted, rotated automatically, and accessed only by authorized IAM roles.

For AI applications, Secrets Manager is commonly used to store API keys for OpenAI, Anthropic, databases, Redis, and other external services.

| WHY INTERVIEWERS ASK THIS

They want to evaluate whether you follow secure development practices.

| COMMON MISTAKES

- Storing API keys inside source code.
- Uploading secrets to GitHub.
- Sharing credentials across multiple applications.

| FOLLOW-UP QUESTIONS

- How is Secrets Manager different from Parameter Store?
- Can Secrets Manager rotate passwords automatically?
- How does IAM work with Secrets Manager?

PRODUCTION INSIGHT

One of the biggest security mistakes developers make is hardcoding API keys. Production systems always store sensitive credentials in a secure secret management service.

What is AWS KMS, and why is encryption important?

What is AWS KMS, and why is encryption important?

| ANSWER

AWS KMS (Key Management Service) is a managed service used to create and manage encryption keys.

Encryption protects sensitive data by converting it into an unreadable format that can only be accessed with the correct key.

Many AWS services such as S3, RDS, EBS, DynamoDB, and Secrets Manager integrate directly with KMS.

For example, if your AI application stores confidential user documents in S3, enabling KMS encryption ensures that the data remains protected even if someone gains unauthorized access to the storage.

KMS simplifies encryption while allowing organizations to control who can use encryption keys.

| WHY INTERVIEWERS ASK THIS

Security is a major part of production AI systems, and interviewers want to know whether you understand data protection.

| COMMON MISTAKES

- Assuming AWS encrypts everything automatically.
- Confusing encryption with authentication.
- Ignoring encryption for stored data.

| FOLLOW-UP QUESTIONS

- What is the difference between encryption at rest and encryption in transit?
- Which AWS services integrate with KMS?
- Can KMS be used with S3?

PRODUCTION INSIGHT

Enterprise AI applications often process sensitive customer data. Encrypting data using AWS KMS is considered a standard security practice and is frequently required for regulatory compliance.

What is AWS CloudFormation, and why is it used?

What is AWS CloudFormation, and why is it used?

| ANSWER

AWS CloudFormation is an Infrastructure as Code (IaC) service that lets you create and manage AWS resources using code instead of manually configuring them through the AWS Console.

Instead of clicking through dozens of settings every time you deploy an application, you write a template that defines your infrastructure. AWS then creates everything automatically.

For example, you can define an EC2 instance, an S3 bucket, an IAM role, a VPC, and a Load Balancer in a single CloudFormation template.

This makes deployments consistent, repeatable, and easier to maintain across development, testing, and production environments.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand modern infrastructure management practices used in production.

| COMMON MISTAKES

- Thinking CloudFormation deploys application code.
- Using manual AWS Console changes in production.
- Confusing CloudFormation with Docker.

| FOLLOW-UP QUESTIONS

- What is Infrastructure as Code?
- How is CloudFormation different from Terraform?
- Can CloudFormation create an entire VPC?

PRODUCTION INSIGHT

Most enterprise teams avoid manually creating AWS resources because manual changes lead to configuration drift. Infrastructure as Code ensures every environment is built the same way.

What is AWS CDK, and how is it different from CloudFormation?

What is AWS CDK, and how is it different from CloudFormation?

| ANSWER

AWS CDK (Cloud Development Kit) lets developers define AWS infrastructure using programming languages like Python, TypeScript, Java, and C#.

Instead of writing lengthy CloudFormation YAML or JSON templates, you write normal code. The CDK automatically converts that code into CloudFormation templates during deployment.

For developers, this is often easier because you can use variables, loops, functions, and reusable components just like any other software project.

CDK is becoming increasingly popular because it combines Infrastructure as Code with modern software engineering practices.

| WHY INTERVIEWERS ASK THIS

They want to know whether you're aware of newer AWS development tools.

| COMMON MISTAKES

- Thinking CDK replaces CloudFormation.
- Assuming CDK directly creates AWS resources.
- Confusing CDK with SDK.

| FOLLOW-UP QUESTIONS

- Which programming languages does CDK support?
- Does CDK use CloudFormation internally?
- Why do developers prefer CDK?

PRODUCTION INSIGHT

Many modern AI teams prefer CDK because developers can manage infrastructure using familiar programming languages instead of manually writing infrastructure templates.

What is Amazon EBS, and how is it different from Amazon S3?

What is Amazon EBS, and how is it different from Amazon S3?

| ANSWER

Amazon EBS (Elastic Block Store) provides block storage for EC2 instances.

Think of EBS as a virtual hard drive attached to a server. It stores operating systems, installed software, application files, and other data required by the EC2 instance.

Amazon S3, on the other hand, is object storage designed to store files like documents, datasets, images, backups, and AI models.

A simple way to remember the difference is:

- EBS stores data for one server.
- S3 stores files that can be accessed independently of any server.

| WHY INTERVIEWERS ASK THIS

They want to test whether you understand different AWS storage services.

| COMMON MISTAKES

- Calling EBS a database.
- Confusing block storage with object storage.
- Thinking EBS can replace S3.

| FOLLOW-UP QUESTIONS

- Can one EBS volume be attached to multiple EC2 instances?
- What happens if an EC2 instance is terminated?
- When should you use S3 instead of EBS?

PRODUCTION INSIGHT

AI applications usually use EBS for application storage and S3 for datasets, trained models, documents, and backups.

What is Amazon EFS, and when should you use it?

What is Amazon EFS, and when should you use it?

| ANSWER

Amazon EFS (Elastic File System) is a managed shared file storage service.

Unlike EBS, which is typically attached to a single EC2 instance, EFS allows multiple EC2 instances to access the same files simultaneously.

This makes it useful when several application servers need access to the same documents, configuration files, or shared datasets.

EFS automatically scales as your storage grows, and AWS manages the underlying infrastructure.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand the different storage options available in AWS.

| COMMON MISTAKES

- Confusing EFS with EBS.
- Assuming EFS is a database.
- Using EFS when object storage is more appropriate.

| FOLLOW-UP QUESTIONS

- What is the difference between EFS and EBS?
- Can multiple EC2 instances access EFS?
- Is EFS managed by AWS?

PRODUCTION INSIGHT

Shared AI applications often use EFS when multiple servers need access to the same files without manually synchronizing data between instances.

How would you deploy a production-ready RAG application on AWS?

How would you deploy a production-ready RAG application on AWS?

| ANSWER

A production RAG application typically uses several AWS services working together.

- Store uploaded documents in Amazon S3.
- Trigger a Lambda function when a new document arrives.
- Generate embeddings using an embedding model.
- Store embeddings in Amazon OpenSearch.
- Deploy the backend API using ECS, EKS, or EC2.
- Use Amazon Bedrock to access the LLM.
- Store chat history in DynamoDB or Redis.
- Monitor logs and system health using CloudWatch.
- Secure access using IAM, Secrets Manager, and KMS.
- Place an Application Load Balancer in front of the API for scalability.

This architecture is scalable, secure, and suitable for real-world enterprise AI applications.

| WHY INTERVIEWERS ASK THIS

They want to evaluate whether you can connect multiple AWS services into a complete production architecture instead of explaining them individually.

| COMMON MISTAKES

- Using only one AWS service for the entire application.
- Forgetting monitoring and security.
- Ignoring scalability and fault tolerance.

| FOLLOW-UP QUESTIONS

- Why use OpenSearch instead of S3 for retrieval?
- Why place a Load Balancer in front of the API?
- Where would you store chat history?

PRODUCTION INSIGHT

Senior AI Engineers are expected to think in terms of complete system architecture rather than individual services. Understanding how AWS services work together is often the difference between passing and failing production-level interviews.

How would you deploy a FastAPI application on AWS?

How would you deploy a FastAPI application on AWS?

| ANSWER

A common production approach is to package the FastAPI application as a Docker container and deploy it on AWS. The typical deployment flow is:

- Build the FastAPI application.
- Create a Docker image.
- Push the image to Amazon ECR.
- Deploy the container using ECS, EKS, or an EC2 instance.
- Place an Application Load Balancer in front of the application.
- Point your domain to the Load Balancer using Route 53.
- Monitor logs and application health using CloudWatch.

This architecture is scalable, easy to maintain, and widely used in production AI applications.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how AI APIs are deployed in real production environments.

| COMMON MISTAKES

- Running FastAPI directly on a local machine.
- Skipping Docker.
- Forgetting monitoring and Load Balancer configuration.

| FOLLOW-UP QUESTIONS

- Why use Docker?
- Can FastAPI run on ECS?
- Why is an Application Load Balancer needed?

PRODUCTION INSIGHT

Most production AI APIs are containerized before deployment because containers make applications portable, consistent, and easier to scale.

How do you secure an AI application running on AWS?

How do you secure an AI application running on AWS?

| ANSWER

Security is one of the most important aspects of deploying AI applications. Some common best practices include:

- Use IAM Roles instead of hardcoded AWS credentials.
- Store API keys in Secrets Manager.
- Encrypt sensitive data using AWS KMS.
- Place databases inside private subnets.
- Restrict network access using Security Groups.
- Enable HTTPS for all APIs.
- Monitor suspicious activity using CloudWatch and CloudTrail.
- Follow the principle of least privilege.

These practices help protect user data, AI models, and cloud resources from unauthorized access.

| WHY INTERVIEWERS ASK THIS

Security is a critical part of production systems, and employers expect AI Engineers to understand basic cloud security.

| COMMON MISTAKES

- Hardcoding API keys.
- Giving AdministratorAccess to every application.
- Exposing databases directly to the internet.

| FOLLOW-UP QUESTIONS

- Why use IAM Roles instead of Access Keys?
- What is the principle of least privilege?
- Where should API keys be stored?

PRODUCTION INSIGHT

Many cloud security incidents happen because of poor IAM permissions or exposed credentials — not because AWS itself is insecure.

What is the difference between EC2, ECS, EKS, Lambda, and Fargate?

What is the difference between EC2, ECS, EKS, Lambda, and Fargate?

| ANSWER

These services all run applications, but they solve different problems.

- EC2 gives you a virtual server that you manage yourself.
- ECS runs Docker containers with AWS-managed orchestration.
- EKS runs Kubernetes clusters for enterprise-scale container management.
- Lambda runs short-lived functions without managing servers.
- Fargate runs Docker containers without requiring you to manage EC2 instances.

The choice depends on your application's requirements. For example, a chatbot API may run on ECS with Fargate, while a document upload trigger may use Lambda.

| WHY INTERVIEWERS ASK THIS

This is one of the most common AWS interview questions because it tests your understanding of AWS compute services.

| COMMON MISTAKES

- Saying one service replaces all the others.
- Confusing ECS with EKS.
- Assuming Lambda is suitable for every workload.

| FOLLOW-UP QUESTIONS

- When would you choose EC2?
- Why use Kubernetes?
- What workloads are best suited for Lambda?

PRODUCTION INSIGHT

Choosing the right compute service affects scalability, operational complexity, performance, and cost. Experienced engineers choose based on the workload rather than personal preference.

How do you monitor an AI application after deployment?

How do you monitor an AI application after deployment?

| ANSWER

Deploying an application is only the beginning. After deployment, continuous monitoring is essential. A production AI application should monitor:

- API response time
- Error rates
- CPU and memory usage
- Request volume
- Application logs
- Model inference latency
- Failed requests
- Database performance

AWS CloudWatch collects these metrics and logs. You can also configure alarms to notify your team whenever something unusual happens. Good monitoring helps detect problems before users report them.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand production operations, not just application development.

| COMMON MISTAKES

- Deploying an application without monitoring.
- Looking only at logs.
- Ignoring application performance metrics.

| FOLLOW-UP QUESTIONS

- What is the difference between metrics and logs?
- Can CloudWatch send alerts?
- What metrics are most important for an AI API?

PRODUCTION INSIGHT

A production AI application is only as good as its monitoring. Fast detection of failures often prevents major outages and improves user experience.

If you were asked to build a scalable AI chatbot on AWS, what architecture wo...

If you were asked to build a scalable AI chatbot on AWS, what architecture would you choose?

| ANSWER

A scalable AI chatbot can be built using several AWS services working together. A common architecture is:

- Route 53 for the domain.
- Application Load Balancer to distribute traffic.
- ECS with Fargate (or EKS) to run the FastAPI backend.
- Amazon Bedrock for LLM inference.
- Amazon OpenSearch for vector search.
- Amazon S3 for storing uploaded documents.
- DynamoDB or Redis for chat history.
- Lambda for background processing.
- SQS for asynchronous tasks.
- CloudWatch for monitoring.
- IAM, Secrets Manager, and KMS for security.

This architecture supports scalability, security, and high availability while keeping the system easy to maintain.

| WHY INTERVIEWERS ASK THIS

They want to evaluate your ability to design an end-to-end production AI system rather than explain individual AWS services.

| COMMON MISTAKES

- Using only one AWS service for everything.
- Ignoring security and monitoring.
- Forgetting asynchronous processing.

| FOLLOW-UP QUESTIONS

- Why choose ECS over EC2?
- Why use SQS in this architecture?
- Where would embeddings be stored?
- How would you scale the chatbot during peak traffic?

PRODUCTION INSIGHT

This type of architecture is commonly used in enterprise GenAI applications because it separates responsibilities across services, making the system more scalable, fault-tolerant, and easier to maintain.

What is the difference between Horizontal Scaling and Vertical Scaling?

What is the difference between Horizontal Scaling and Vertical Scaling?

| ANSWER

Scaling means increasing the capacity of your application to handle more users or requests.

Vertical Scaling means increasing the resources of a single server, such as adding more CPU, RAM, or GPU. For example: upgrade an EC2 instance from 4 vCPUs to 16 vCPUs.

Horizontal Scaling means adding more servers instead of making one server bigger. For example: increase your application from 2 ECS containers to 10 containers during high traffic.

Horizontal scaling is generally preferred in production because it improves availability and allows the application to continue running even if one server fails.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand how production systems handle increasing workloads.

| COMMON MISTAKES

- Thinking scaling always means upgrading one server.
- Ignoring fault tolerance.
- Confusing Auto Scaling with Vertical Scaling.

| FOLLOW-UP QUESTIONS

- Which scaling method is more common in AWS?
- Can ECS scale horizontally?
- Which AWS service performs Auto Scaling?

PRODUCTION INSIGHT

Modern AI applications almost always use horizontal scaling because it's easier to handle traffic spikes and provides better fault tolerance than relying on a single powerful server.

What is High Availability (HA), and why is it important?

What is High Availability (HA), and why is it important?

| ANSWER

High Availability means designing an application so it continues running even if one component fails.

Instead of depending on a single EC2 instance or database, production systems deploy multiple resources across different Availability Zones. For example: multiple ECS tasks, multiple EC2 instances, Load Balancer, Multi-AZ RDS, Auto Scaling.

If one server crashes, another server immediately starts handling requests, so users experience little or no downtime.

High Availability is essential for production AI applications that serve customers 24×7.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand reliability and fault tolerance.

| COMMON MISTAKES

- Deploying everything on one EC2 instance.
- Assuming backups provide High Availability.
- Ignoring Availability Zones.

| FOLLOW-UP QUESTIONS

- What is an Availability Zone?
- What is the difference between HA and Disaster Recovery?
- How does a Load Balancer improve availability?

PRODUCTION INSIGHT

Enterprise AI systems are designed assuming failures will happen. High Availability ensures the application keeps running even when servers fail.

What is the difference between an Availability Zone and an AWS Region?

What is the difference between an Availability Zone and an AWS Region?

| ANSWER

An AWS Region is a geographical location where AWS has data centers. Examples include Mumbai, Singapore, Frankfurt, and Virginia.

Each Region contains multiple Availability Zones (AZs). An Availability Zone is an isolated data center with independent power, cooling, and networking.

For example, the Mumbai Region contains multiple Availability Zones.

If one Availability Zone experiences a failure, applications deployed across multiple AZs can continue serving users without interruption.

| WHY INTERVIEWERS ASK THIS

This is one of the most frequently asked AWS fundamentals in cloud interviews.

| COMMON MISTAKES

- Thinking a Region contains only one data center.
- Confusing Region with Availability Zone.
- Deploying everything in a single AZ.

| FOLLOW-UP QUESTIONS

- Why deploy applications across multiple AZs?
- Can an ALB distribute traffic across AZs?
- Does every Region have multiple Availability Zones?

PRODUCTION INSIGHT

Most production workloads are deployed across multiple Availability Zones to improve availability and reduce downtime during infrastructure failures.

What happens when an EC2 instance suddenly crashes?

What happens when an EC2 instance suddenly crashes?

| ANSWER

If an EC2 instance crashes, the impact depends on how the application is designed.

In a production environment, the Application Load Balancer detects that the instance has become unhealthy and automatically stops sending traffic to it.

If Auto Scaling is enabled, AWS launches a new EC2 instance to replace the failed one.

If the application is running across multiple instances, users may not even notice that one server has failed.

Without Auto Scaling or a Load Balancer, however, the application could become unavailable until someone manually fixes the issue.

| WHY INTERVIEWERS ASK THIS

They want to evaluate whether you understand fault tolerance in cloud environments.

| COMMON MISTAKES

- Thinking AWS automatically recovers every EC2 instance.
- Ignoring Load Balancers.
- Assuming one server is enough for production.

| FOLLOW-UP QUESTIONS

- How does Auto Scaling replace failed instances?
- Who detects unhealthy servers?
- What happens if Auto Scaling is disabled?

PRODUCTION INSIGHT

Production systems are designed so that server failures are expected events rather than emergencies. Automation replaces failed resources without requiring manual intervention.

What AWS services have you used in your AI projects?

What AWS services have you used in your AI projects?

| ANSWER

A good interview answer should be based on your actual experience. For example:

"In my RAG project, I used Amazon S3 to store uploaded PDF documents. The backend API was built with FastAPI and deployed as a Docker container. I stored the container image in Amazon ECR and deployed it using ECS. IAM Roles handled permissions, while Secrets Manager securely stored API keys. OpenSearch was used for vector search, CloudWatch monitored logs and performance, and Bedrock provided access to the LLM. This architecture was scalable, secure, and easy to maintain."

This answer demonstrates that you understand how AWS services work together in a real application.

| WHY INTERVIEWERS ASK THIS

They want to assess whether you've actually built AI systems using AWS rather than simply memorizing service definitions.

| COMMON MISTAKES

- Listing AWS services without explaining their purpose.
- Mentioning services you've never used.
- Giving theoretical answers without describing a workflow.

| FOLLOW-UP QUESTIONS

- Why did you choose ECS instead of EC2?
- Why store documents in S3?
- What monitoring tools did you use?
- How did you secure your application?

PRODUCTION INSIGHT

Interviewers are usually more interested in how you connected AWS services together than in how many AWS services you can name. A clear explanation of your architecture often leaves a stronger impression than memorizing definitions.

How do you optimize AWS costs for an AI application?

How do you optimize AWS costs for an AI application?

| ANSWER

Managing cloud costs is just as important as building the application itself. A well-designed AI system should be scalable without wasting resources. Some common ways to reduce AWS costs are:

- Stop unused EC2 instances.
- Use Auto Scaling so servers run only when needed.
- Store infrequently accessed data in cheaper S3 storage classes.
- Use Spot Instances for non-critical workloads like model training.
- Delete unused EBS volumes and snapshots.
- Right-size EC2 instances instead of choosing oversized machines.
- Monitor costs using AWS Cost Explorer and AWS Budgets.

The goal is to pay only for the resources your application actually needs.

| WHY INTERVIEWERS ASK THIS

Cloud costs can become very high for AI applications. Interviewers want to know whether you can build solutions that are both scalable and cost-effective.

| COMMON MISTAKES

- Choosing the largest EC2 instance without checking actual usage.
- Leaving GPU instances running after work is finished.
- Ignoring AWS cost monitoring tools.

| FOLLOW-UP QUESTIONS

- What are Spot Instances?
- How does Auto Scaling reduce costs?
- Which AWS service helps monitor cloud spending?

PRODUCTION INSIGHT

Many AI startups spend more money on cloud infrastructure than necessary. Engineers who understand cost optimization are highly valued because they directly reduce operational expenses.

What are Spot Instances, and when should you use them?

What are Spot Instances, and when should you use them?

| ANSWER

Spot Instances are unused EC2 instances that AWS offers at significantly lower prices than On-Demand instances.

They are a great option for workloads that can tolerate interruptions because AWS may reclaim the instance when it needs the capacity.

AI Engineers commonly use Spot Instances for model training, data preprocessing, batch inference, embedding generation, and large offline AI jobs.

However, they are generally not recommended for production APIs that require continuous availability.

| WHY INTERVIEWERS ASK THIS

They want to know whether you understand AWS pricing models and cost optimization strategies.

| COMMON MISTAKES

- Running production APIs entirely on Spot Instances.
- Assuming Spot Instances never stop unexpectedly.
- Confusing Spot Instances with Reserved Instances.

| FOLLOW-UP QUESTIONS

- What happens if AWS reclaims a Spot Instance?
- When should you avoid using Spot Instances?
- How much cheaper are Spot Instances?

PRODUCTION INSIGHT

Many companies use Spot Instances for ML training because training jobs can usually resume from checkpoints, making interruptions less problematic while saving substantial costs.

Explain a complete AWS architecture for an AI document chatbot.

Explain a complete AWS architecture for an AI document chatbot.

| ANSWER

A production AI document chatbot typically follows this architecture:

- Users upload PDF files to Amazon S3.
- An S3 event triggers an AWS Lambda function.
- Lambda extracts the document text.
- An embedding model generates embeddings.
- Embeddings are stored in Amazon OpenSearch.
- The FastAPI backend runs on ECS with Fargate.
- Amazon Bedrock generates responses using the retrieved context.
- Chat history is stored in DynamoDB or Redis.
- CloudWatch monitors logs and application health.
- IAM, Secrets Manager, and KMS secure the entire application.

This architecture is scalable, secure, and follows AWS best practices for production AI systems.

| WHY INTERVIEWERS ASK THIS

This is a common system design question that evaluates whether you can connect multiple AWS services into a complete AI solution.

| COMMON MISTAKES

- Forgetting document storage.
- Ignoring monitoring and security.
- Using the LLM without a retrieval layer.

| FOLLOW-UP QUESTIONS

- Why use OpenSearch?
- Why trigger Lambda from S3?
- Where would you store conversation history?
- How would you scale this application?

PRODUCTION INSIGHT

Strong candidates explain not only which AWS services they use, but also why each service is responsible for a specific part of the architecture.

If your AI API suddenly becomes very slow, how would you troubleshoot it?

If your AI API suddenly becomes very slow, how would you troubleshoot it?

| ANSWER

I would troubleshoot the issue step by step instead of making assumptions. My approach would be:

- Check CloudWatch logs for errors.
- Monitor CPU and memory usage.
- Check API response times.
- Verify database performance.
- Check Bedrock or LLM response latency.
- Inspect OpenSearch query performance.
- Verify whether Auto Scaling has added enough instances.
- Check network latency and recent deployments.
- Identify whether the issue affects all users or only specific requests.

By isolating each component, it's easier to identify the real bottleneck.

| WHY INTERVIEWERS ASK THIS

They want to evaluate your production debugging skills rather than your theoretical AWS knowledge.

| COMMON MISTAKES

- Restarting servers without investigation.
- Assuming the LLM is always the problem.
- Ignoring monitoring dashboards.

| FOLLOW-UP QUESTIONS

- Which CloudWatch metrics would you check first?
- How would you identify whether the database is the bottleneck?
- How would you reduce API latency?

PRODUCTION INSIGHT

Good AI Engineers troubleshoot using data and monitoring tools instead of guessing. A structured debugging process often matters more than finding the answer immediately.

What are the most important AWS services every AI Engineer should know for in...

What are the most important AWS services every AI Engineer should know for interviews?

| ANSWER

For most AI Engineer interviews, these AWS services appear repeatedly:

IAM EC2 S3 Lambda API Gateway ECS EKS Fargate ECR Bedrock
SageMaker OpenSearch DynamoDB RDS SQS SNS EventBridge Step Functions
CloudWatch Route 53 CloudFront Auto Scaling Load Balancer Secrets Manager KMS
CloudFormation CDK

Rather than memorizing definitions, understand why each service exists, when to use it, and how they work together in a production AI architecture.

| WHY INTERVIEWERS ASK THIS

Senior interviewers care more about your ability to design complete cloud solutions than your ability to recite AWS documentation.

| COMMON MISTAKES

- Memorizing services without understanding their purpose.
- Learning services in isolation.
- Ignoring security, monitoring, and deployment topics.

| FOLLOW-UP QUESTIONS

- Which AWS services have you used together?
- Design an AI architecture using these services.
- Which AWS services would you choose for a GenAI chatbot?
- Which services are most important for RAG?

PRODUCTION INSIGHT

If you can confidently explain how services like S3 → Lambda → OpenSearch → Bedrock → ECS → CloudWatch → IAM work together in one production architecture, you'll answer the majority of AWS interview questions for AI Engineer roles with confidence.



YOU FINISHED THE FREE GUIDE

You now know the **AWS** fundamentals interviewers test for.

You've completed all 50 questions covering compute, storage, networking, GenAI services, security, and production architecture on AWS.

50

Q&A Completed

27

AWS Services Covered

4

Difficulty Levels

PREMIUM EDITION

Go deeper with 200 Q&A — Production Scenarios & System Design

The free guide covered the fundamentals. The Premium Edition goes much further into what actually gets asked at the Senior and Principal level.

- ✓ Advanced Bedrock & SageMaker production workflows
- ✓ Kubernetes, GPU workloads & multi-region architectures
- ✓ CI/CD, disaster recovery & real debugging scenarios
- ✓ Security, networking & deployment deep-dives

Get the Premium Edition →

roy-s-ai-lab.vercel.app



Ritesh Rai (Roy)
GenAI Engineer · Founder, Roy's AI Lab